


# Keras Core

Back to being a polyglot of Deep-Learning



# Hi!

I am Dilip Parasu.

Also known as SuperSecureHuman



<https://github.com/SuperSecureHuman>



<https://linkedin.com/in/dilip-parasu>



<https://supersecurehuman.github.io/>

# Good old Keras

**Once upon a time** there existed Keras

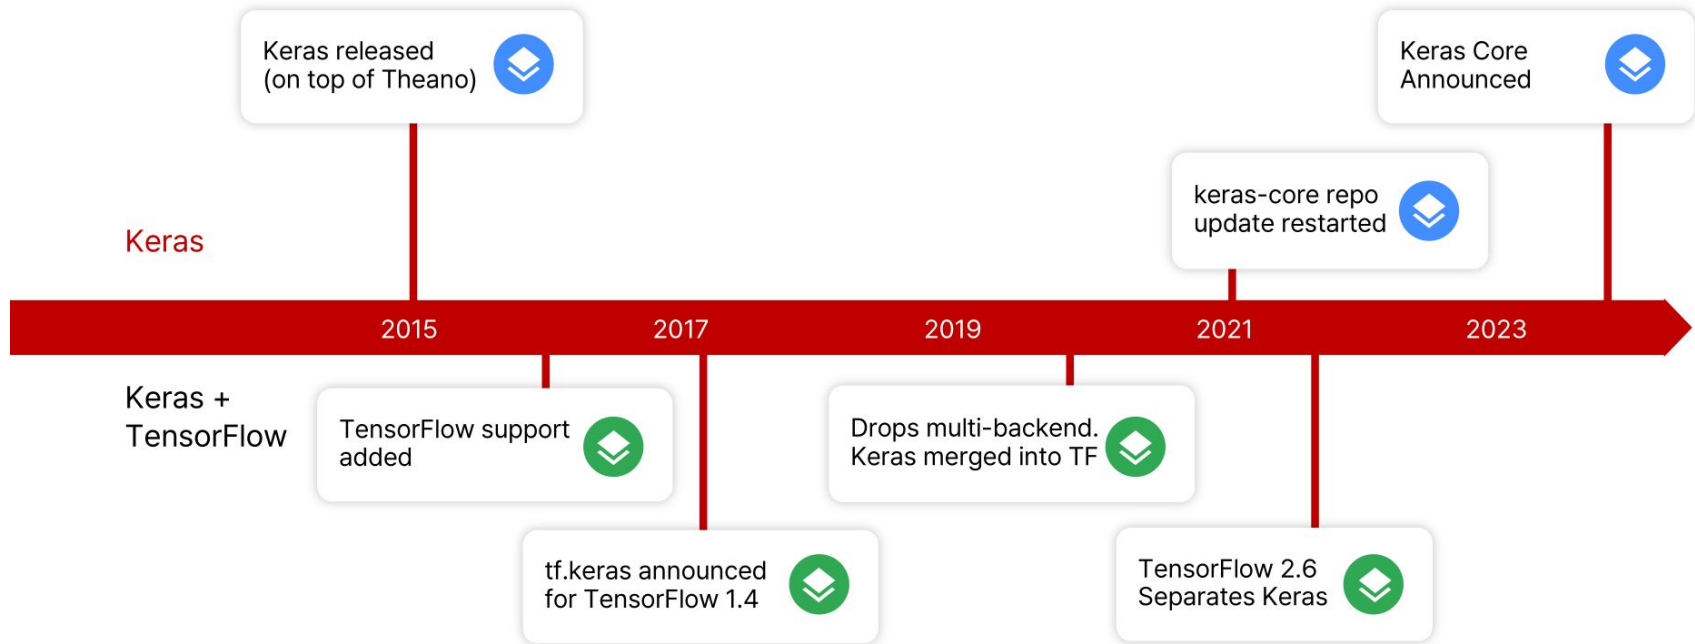
- High performance Deep learning Library for python
- That was very easy to use
- Supported Theano, MXNET, TFv1, CNTK as backend

**Then**

- It became part of Tensorflow
- Dropped other backend support
- Name space was in tf.keras




# From Multi to Single to Multi



# Now where are we?

- Switchable backend (TF, JAX, Pytorch)
- Multi framework custom components
- Universal training loop
- Native models support
- Future proof code



```
model = get_keras_core_model()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
loss_fn = torch.nn.CrossEntropyLoss()

def train_step(inputs, targets):
    # Compute loss.
    logits = model(inputs, training=True)
    loss = loss_fn(logits, targets)

    # Compute gradients.
    model.zero_grad()
    loss.backward()

    # Update weights.
    optimizer.step()
    return loss

# Iterate over epochs.
for epoch in range(num_epochs):
    # Iterate over the batches of the dataset.
    for step, (inputs, targets) in enumerate(dataset):
        loss = train_step(inputs, targets)
        print(f'Loss: {loss.detach().numpy():.4f}')
```



```
model = get_keras_core_model()
optimizer = keras.optimizers.Adam(learning_rate=1e-3)
loss_fn = keras.losses.CategoricalCrossentropy(from_logits=True)

@tf.function(jit_compile=True)
def train_step(inputs, targets):
    # Compute loss.
    with tf.GradientTape() as tape:
        logits = model(inputs, training=True)
        loss = loss_fn(targets, logits)

    # Compute gradients.
    gradients = tape.gradient(loss, model.trainable_weights)

    # Update weights.
    optimizer.apply_gradients(model.trainable_weights)
    return loss

# Iterate over epochs.
for epoch in range(num_epochs):
    # Iterate over the batches of the dataset.
    for step, (inputs, targets) in enumerate(dataset):
        loss = train_step(inputs, targets)
        print(f'Loss: {loss.numpy():.4f}')
```

Writing a custom  
training loop for  
a Keras model:  
PyTorch,  
TensorFlow,  
JAX



```
model = get_keras_core_model()
optimizer = keras.optimizers.Adam(learning_rate=1e-3)
loss_fn = keras.losses.CategoricalCrossentropy(from_logits=True)

# All variables must be built before training starts.
optimizer.build(model.trainable_variables)

def compute_loss_and_updates(trainable_vars, non_trainable_vars, data):
    # Stateless function to compute the loss and non-trainable var updates.
    x, y = data
    pred, non_trainable_vars = model.stateless_call(trainable_vars, non_trainable_vars, x)
    loss = loss_fn(y, pred)
    return loss, non_trainable_vars

# Function that returns the gradients for the trainable vars.
grad_fn = jax.value_and_grad(compute_loss_and_updates, has_aux=True)

@jax.jit
def train_step(state, data):
    # Stateless function that calls the grad_fn and computes trainable vars updates.
    trainable_vars, non_trainable_vars, optimizer_vars = state
    print(len(trainable_vars), len(non_trainable_vars), len(optimizer_vars))
    (loss, non_trainable_vars), grads = grad_fn(trainable_vars, non_trainable_vars, data)
    trainable_vars, optimizer_vars = optimizer.stateless_apply(
        optimizer_vars, grads, trainable_vars
    )
    return loss, (trainable_vars, non_trainable_vars, optimizer_vars)

# Prepare model state.
state = (model.trainable_variables, model.non_trainable_variables, optimizer.variables)

# Iterate over epochs.
for epoch in range(num_epochs):
    # Iterate over the batches of the dataset.
    for step, (inputs, targets) in enumerate(dataset):
        # Each train_step call is entirely stateless (no side effects).
        loss, state = train_step(state, data)
        print(f'Loss: {loss:.4f}')
```

# Switchable backend

2 Lines, and you are good to go!

Why is this even a deal?

2 big players - Tensorflow and Pytorch with having almost 40 to 60% of share each.

Imagine now having just one codebase, and you have covered almost 100% of the users.

Sounds like flutter → IOS and Android, Single Codebase



```
import os
os.environ["KERAS_BACKEND"] = 'jax'
# Set the backend first always

import keras_core as keras
```



# Seamless integration with backends

Train Keras models with

- Low level JAX - `optax`, `jax.grad`, `jax.jit`, `jax.pmap` ...
- Low level tf - `tf.GradientTape`...
- Low level torch - `torch.optim`, `torch.nn`

You can use the keras model you made, directly in torch as a `nn.Module` !




# JAX enters

**Blazing Speed:** JAX delivers high-performance computing with GPU/TPU acceleration, XLA compilation, perfect for large-scale numerical tasks.

**Effortless Gradients:** JAX's automatic differentiation simplifies gradient computations.

**Functional Magic:** JAX promotes functional programming, enabling modular and readable code for machine learning pipelines.

**Scalable Magic:** JAX streamlines parallelism and distributed computing (via vmap, pmap), ideal for scaling computations across multiple devices.





# What powers JAX



## **XLA (Accelerated Linear Algebra)**

- Fast matrix operations on CPU, GPU, TPU
- JIT!



## **MLIR (Multi Level Intermediate Representation)**

- Intermediate compilation layer for ML



## **OpenXLA**

- Open source version of XLA and StableHLO
  - StableHLO - Operation set for high level operations HLO in ML models



# What is around JAX



- The core



## FLAX

- Based on Jax
- High level API for defining and Training Neural Networks



## Haiku

- Another high level API



## Optax

- Optimizers and Loss functions



## Huggingface

- Probably the largest model source for jax now?



# So what's the catch

- Super high learning curve
- Still growing, and changing docs
- Lack of 3rd party codes
- You need to care about distributed workload
- Purely Stateless



# Keras helps here

- You don't have to care about the stateless nature of jax
- All the performance optimizations taken care (up to a point)
- Using JAX while maintaining a very familiar, easy to read implementation



Introducing Keras Core:  
Keras for TensorFlow, JAX, and PyTorch.





# To the Notebooks

